

Project Title: Kubernetes Cluster Security Assessment

Short Project Description:

The system generates actionable reports to improve cluster security posture.

Component	Role
Cluster Inventory Collector	Gathers information on nodes, pods, services
Role-Based Access Control (RBAC) Analyzer	Analyzes Kubernetes roles and permissions
Pod Security Evaluator	Reviews pod security settings
Network Policy Analyzer	Evaluates inter-pod network segmentation
Report Generator	Creates cluster security assessment reports

Component Details:

- Cluster Inventory Collector:**
 - Connects to Kubernetes API.
 - Collects data on:
 - Nodes
 - Pods
 - Namespaces
 - Etc
 - RBAC Analyzer:**
 - Looks for:
 - Overly permissive roles
 - Missing namespace separation
 - Etc
 - Pod Security Evaluator:**
 - Checks:
 - If pods run as root
 - Lack of security contexts
 - Etc
 - Network Policy Analyzer:**
 - Verifies if network policies limit traffic between pods.
 - Report Generator:**
 - Creates prioritized action items to secure the cluster.
-

Overall System Flow:

- Input: Access to Kubernetes API
- Output: Kubernetes cluster security audit report
- Focus on **configuration, network, and identity management auditing**

Internal Functioning of Each Module:

1. Cluster Inventory Collector

- **How it works:**
 - Authenticates with kubeconfig or in-cluster API credentials.
 - Enumerates:
 - Nodes
 - Pods
 - Services
 - Ingress controllers
 - Network policies
 - Etc
 - Pulls metadata for each object (labels, annotations, service accounts).

2. RBAC Analyzer

- **How it works:**
 - Parses RoleBindings, ClusterRoleBindings, Roles, and ClusterRoles.
 - Checks for:
 - Privileged roles (admin, cluster-admin)
 - Wildcard permissions ("verbs: *", "resources: *")
 - Roles granted outside intended namespaces
 - Etc
- **Flags:** Over-privileged service accounts and risky role grants.

3. Pod Security Evaluator

- **How it works:**
 - Parses pod specs.
 - Checks for:
 - runAsUser=0 (root pods)
 - hostPID, hostNetwork access
 - Missing securityContext or seccompProfile
 - Privileged container flags (privileged: true)
 - Etc

4. Network Policy Analyzer

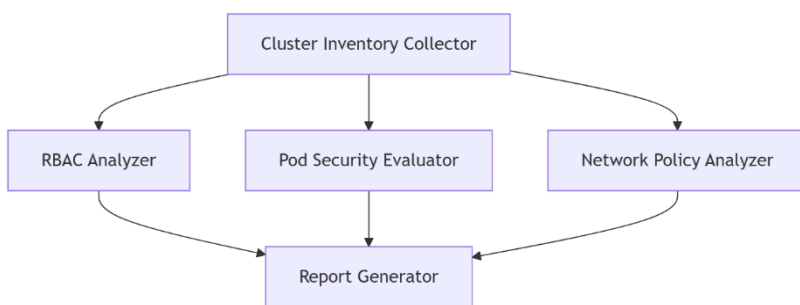
- **How it works:**
 - Lists all NetworkPolicy objects.
 - Verifies if:

- Segmentation exists between namespaces.
 - Traffic is restricted (deny-all first policy).
 - **Flags:** Open communication between sensitive pods.
-

5. Report Generator

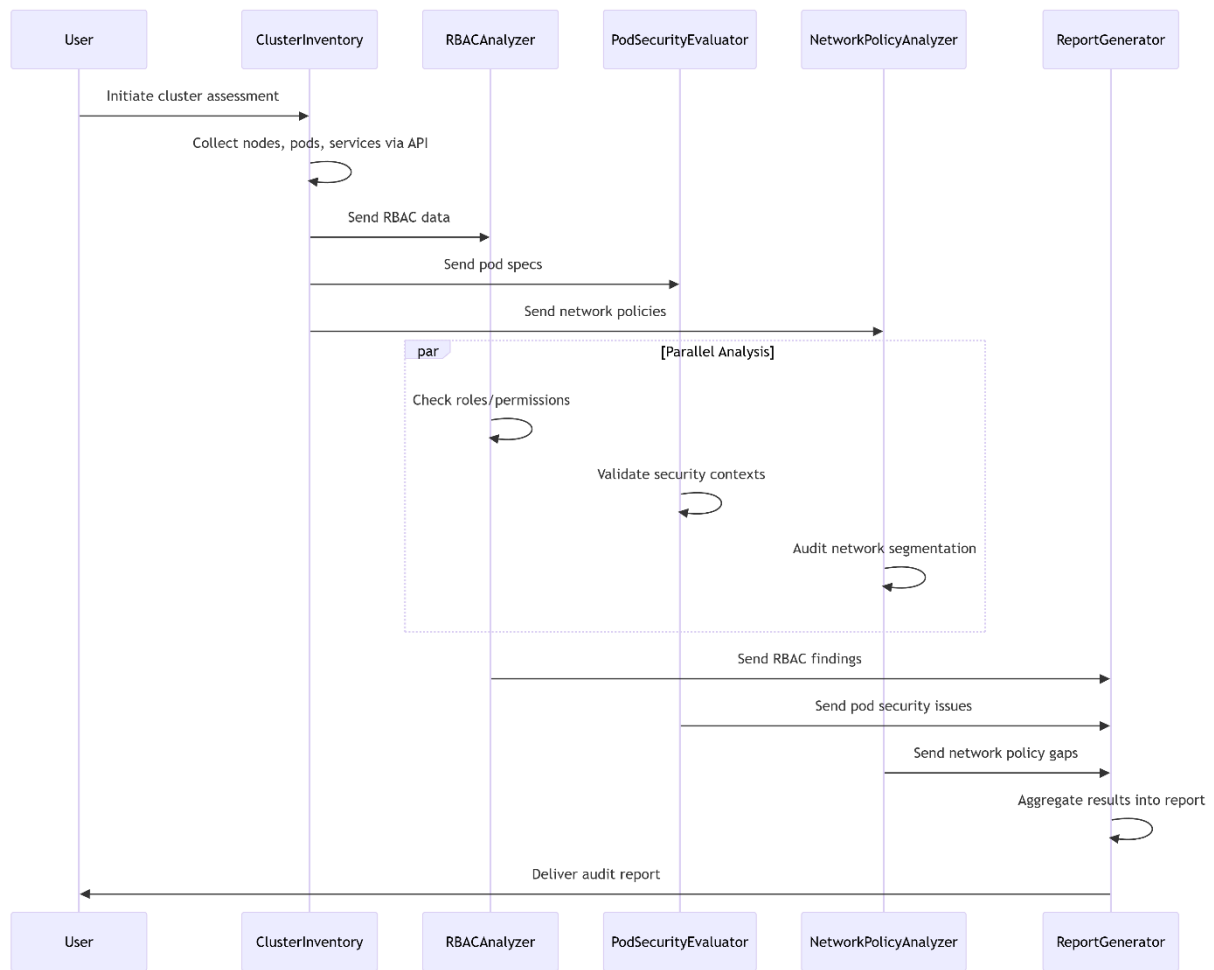
- **How it works:**
 - Aggregates results into:
 - Configuration risks
 - Identity management risks
 - Network segmentation gaps
 - Suggests best practice remediations (e.g., implement deny-all network policy).
-

Component Diagram



- The **Cluster Inventory Collector** gathers cluster metadata and distributes it to three analyzers: **RBAC Analyzer**, **Pod Security Evaluator**, and **Network Policy Analyzer**.
- All analyzers send their findings to the **Report Generator**, which consolidates the results into a structured report.

Sequence Diagram



- The **User** triggers the assessment, prompting the **Cluster Inventory Collector** to fetch data from the Kubernetes API.
- The inventory data is sent to the analyzers for parallel processing:
 - **RBAC Analyzer** checks role-based access controls.
 - **Pod Security Evaluator** reviews pod configurations.
 - **Network Policy Analyzer** audits network segmentation rules.
- Results are aggregated by the **Report Generator**, which produces a prioritized security report for the **User**.

Detailed Project Description: Kubernetes Cluster Security Assessment

A tool for auditing Kubernetes clusters by analyzing configurations, access controls, pod security, and network policies. The tool generates actionable reports to improve cluster security posture.

1. System Overview

The tool assesses Kubernetes clusters for common security risks, including:

- **RBAC Misconfigurations:** Overly permissive roles and service accounts.
 - **Pod Security Gaps:** Containers running as root, privileged mode.
 - **Network Segmentation Issues:** Lack of network policies or open pod-to-pod traffic.
 - **Etc.**
-

2. Component Design & Implementation

2.1 Cluster Inventory Collector

Functionality:

- Gathers cluster metadata (nodes, pods, services) via the Kubernetes API.

Implementation Steps (e.g.):

1. Authentication:

- Use `kubeconfig` files or in-cluster service accounts for API access.

2. Data Collection:

- Use Kubernetes client libraries (`kubectl`, `client-go` for Go, `kubernetes` for Python) to fetch:
 - Nodes: `kubectl get nodes -o json`
 - Pods: `kubectl get pods --all-namespaces -o json`

- Network Policies: `kubectl get networkpolicies -o json`

3. Data Structuring:

- Store collected data in JSON format for downstream analysis.

Output:

- Structured inventory of cluster resources.
-

2.2 RBAC Analyzer

Functionality:

- Identifies overly permissive roles and role bindings.

Implementation Steps (e.g.):

1. Data Parsing:

- Extract Roles, ClusterRoles, RoleBindings, and ServiceAccounts from inventory.

2. Rule Engine:

- Define YAML rules for risky configurations
- Example

```
rules:
  - id: "WILDCARD_VERBS"
    description: "Role allows wildcard verbs (*)"
    condition: "verbs: '*'"
    severity: "High"
```

- Check for:
 - Wildcards (`resources: "*" , verbs: "*" ,`)
 - Cluster-admin roles assigned to non-system accounts,
 - Roles granted across namespaces.

3. Tools:

- Reference open-source tools like **kubectl-who-can** or **rbac-lookup**.

Output:

- List of RBAC violations with severity levels.
-

2.3 Pod Security Evaluator

Functionality:

- Audits pods for insecure configurations.

Implementation Steps (e.g.):

1. Pod Spec Analysis:

- Parse `securityContext` and `pod.spec` fields.

2. Security Checks:

- Flag pods with:
 - `runAsUser: 0` (root execution),
 - `privileged: true`,
 - Missing `securityContext` or `seccompProfile`,
 - Use of `hostNetwork` or `hostPID`.

3. Benchmarking:

- Align checks with **Kubernetes Pod Security Standards** (Baseline/Restricted).

Output:

- List of non-compliant pods and recommended fixes.
-

2.4 Network Policy Analyzer

Functionality:

- Evaluates network segmentation and policy enforcement.

Implementation Steps (e.g.):

1. Policy Extraction:

- Fetch existing NetworkPolicy objects.

2. Rule Checks:

- Verify if namespaces have a default "deny-all" policy.
- Identify pods exposed to all traffic (`ingress: {}`).
- Check for overly permissive rules (e.g., `cidr: 0.0.0.0/0`).

3. Tools:

- Use **Cilium** or **Calico** for policy visualization.

Output:

- List of network exposure risks and policy gaps.
-

2.5 Report Generator

Functionality:

- Consolidates findings into prioritized reports.

Implementation Steps (e.g.):

1. Data Aggregation:

- Merge results from RBAC, Pod Security, and Network Policy analyzers.

2. Report Templates:

- **HTML**: Use Jinja2 for interactive tables and filters.
- **PDF**: Generate via `WeasyPrint` or `pandoc`.
- **SARIF**: For CI/CD integration (e.g., GitHub Advanced Security).
- **Etc.**

3. Prioritization:

- Rank findings by severity (Critical > High > Medium > Low).
- Include remediation steps (e.g., "Restrict role verbs to specific actions").

Output:

- Final report with:

- RBAC violations,
 - Pod security misconfigurations,
 - Network policy gaps,
 - Step-by-step remediation guidance.
-

3. Technology Stack (e.g.)

- **Kubernetes Interaction:** `kubectl`, `client-go`, `kubernetes` Python library, etc.
 - **RBAC Analysis:** Custom rule engine, `rbac-police`, etc.
 - **Pod Security:** Open Policy Agent (OPA) for policy checks, etc.
 - **Reporting:** Jinja2, WeasyPrint, SARIF SDK, etc.
-

4. Evaluation & Validation

1. **Test Clusters:**
 - Deploy clusters with intentional misconfigurations (e.g., privileged pods, wildcard roles, etc).
 2. **Accuracy:**
 - Compare results against manual audits or tools like **kube-bench**.
 3. **Performance:**
 - Measure scan time for clusters with 1000+ pods.
-

5. Development Roadmap

1. **Phase 1:** Build Cluster Inventory Collector and RBAC Analyzer.
2. **Phase 2:** Implement Pod Security and Network Policy analyzers.
3. **Phase 3:** Develop reporting engine with SARIF/HTML support.
4. **Phase 4 (optional):** Validate on real-world clusters (e.g., EKS, GKE).

6. Challenges & Mitigations (optional)

- **Scale:** Optimize API calls with pagination and caching.
- **Complex Policies:** Use OPA or Kyverno for advanced policy validation.
- **False Positives:** Refine rules using community benchmarks like CIS Kubernetes.

7. Glossary

- **RBAC:** Role-Based Access Control
 - **CIDR:** Classless Inter-Domain Routing
 - **SARIF:** Static Analysis Results Interchange Format
 - **OPA:** Open Policy Agent
-