

# Project Title: Contextual Vulnerability Risk Scoring via LLM Reasoning

---

## Core Idea:

- Use LLMs to **analyze each vulnerability** in context (asset type, exposure, business impact) and dynamically assign a **risk score** beyond simple CVSS metrics.
- 

Component	Role
Asset Context Collector	Gathers asset metadata (type, criticality, exposure level, etc)
LLM-Based Risk Reasoning Engine	Reason about impact, exposure, and criticality for each vulnerability
Business Impact Modeler	Calculates operational, financial, reputational risk
Prioritization Score Calculator	Generates dynamic, context-aware risk scores
Reporting Engine	Creates prioritized vulnerability remediation lists

---

## Component Details:

1. **Asset Context Collector:**
  - Pulls:
    - CMDB (Configuration Management Database) data
    - Cloud inventory APIs
    - Asset criticality tags
    - Etc
2. **LLM-Based Risk Reasoning Engine:**
  - Input:
    - Vulnerability description
    - Asset context
  - Output:
    - Dynamic qualitative and quantitative risk reasoning.
3. **Business Impact Modeler:**
  - Simulates:
    - Downtime costs
    - Data breach costs
    - Regulatory fines (e.g., GDPR, HIPAA, NIS 2)
4. **Prioritization Score Calculator:**
  - Combines:
    - LLM qualitative output
    - Business impact models
  - Calculates **dynamic risk scores** far beyond static CVSS.
5. **Reporting Engine:**

- Creates dashboards and reports:
    - Sorted by business risk, not just technical severity.
- 

## Overall System Flow:

- Input: Vulnerability + Asset context
  - Output: Ranked vulnerability list by real-world impact
  - Focus: **Business-aware risk prioritization**
- 

## Internal Functioning of Each Module:

### 1. Asset Context Collector

- **Data sources:**
    - CMDB (asset metadata)
    - Cloud inventory APIs (AWS, Azure Resource Graph, etc)
    - Service type detection (web servers, DBs, control systems, etc)
    - Etc
  - **Metadata collected:**
    - Asset type (server, workstation, IoT, etc)
    - Criticality (high, medium, low business value, etc)
    - Exposure level (public internet, internal network, etc)
    - Etc
- 

### 2. LLM-Based Risk Reasoning Engine

- **Process:**
    - Input:
      - Vulnerability details (CVE, attack vector, complexity etc)
      - Asset context
    - Reasoning using chain-of-thought prompting :
      - « This CVE allows remote unauthenticated access. The target is a public-facing web server hosting customer data → High risk. »
    - Output:
      - Narrative risk explanation.
- 

### 3. Business Impact Modeler

- **Model elements:**
  - Direct costs (downtime × cost/hour)
  - Regulatory fines (HIPAA, NIS 2, GDPR calculations)
  - Brand damage (proxy via historical breach studies)

---

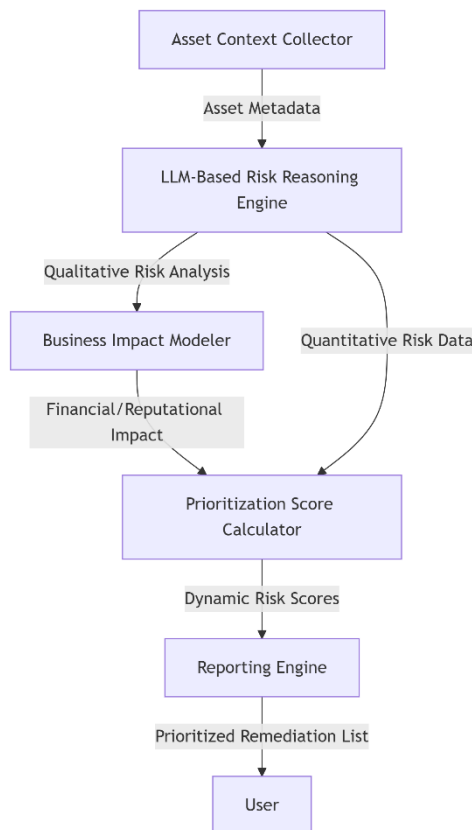
## 4. Prioritization Score Calculator

- **Risk aggregation:**
    - Weighted scoring:
      - $(\text{LLM threat level} \times \text{Asset exposure} \times \text{Business impact})$
  - **Output:**
    - Prioritized list with transparent scoring breakdown.
- 

## 6. Reporting Engine

- **Deliverables:**
    - Graphical risk rankings.
    - Top vulnerabilities.
    - Suggested remediation orders.
    - Etc.
- 


## Component Diagram



1. **Asset Context Collector:** Gathers metadata (e.g., asset type, exposure level, etc) from CMDB, cloud APIs, and criticality tags.

2. **LLM-Based Risk Reasoning Engine:** Analyzes vulnerabilities in context (e.g., "Public web server with customer data → high risk").
3. **Business Impact Modeler:** Estimates operational/financial impact (e.g., downtime costs, GDPR/NIS 2 fines, etc).
4. **Prioritization Score Calculator:** Combines LLM insights and business impact to compute dynamic scores.
5. **Reporting Engine:** Ranks vulnerabilities by business-critical risk, not just CVSS.

## Sequence Diagram



1. **User** submits a vulnerability and asset data (e.g., CVE-2023-XXXX + public-facing server details).
2. **Asset Context Collector** forwards metadata (e.g., "public internet exposure, high criticality") to the **LLM Engine**.
3. **LLM Engine** performs risk reasoning (e.g., "Remote code execution on a customer portal → severe reputational risk").
4. **Business Impact Modeler** calculates financial/reputational costs and shares them with the **Score Calculator**.
5. **Score Calculator** combines LLM outputs and business impact to generate a dynamic risk score (e.g., 9.8/10).
6. **Reporting Engine** delivers a prioritized list (e.g., "Patch this server first due to high business impact").

## Detailed Project Description: Contextual Vulnerability Risk Scoring via LLM Reasoning

A framework that dynamically prioritizes vulnerabilities by combining Large Language Model (LLM) reasoning, asset context, and business impact analysis. This framework moves beyond static CVSS scores to deliver actionable, business-aware risk assessments.

---

### 1. System Architecture Overview

#### Core Components & Interactions

1. **Asset Context Collector**

- *Inputs:* CMDB data, cloud APIs (AWS/Azure), criticality tags, etc.
- *Outputs:* Asset metadata (type, exposure, business criticality, etc).

2. **LLM-Based Risk Reasoning Engine**

- *Inputs:* Vulnerability details (CVE, exploitability, etc), asset context.
- *Outputs:* Narrative risk explanations (e.g., "RCE on public-facing server").

3. **Business Impact Modeler**

- *Inputs:* Asset criticality, industry-specific risk factors.
- *Outputs:* Financial/reputational impact estimates (downtime costs, GDPR / NIS 2, etc, fines).

4. **Prioritization Score Calculator**

- *Inputs:* LLM reasoning, business impact data.
- *Outputs:* Dynamic risk scores (e.g., 9.8/10).

5. **Reporting Engine**

- *Inputs:* Risk scores.
- *Outputs:* Prioritized remediation lists, dashboards, etc.

#### Integration Flow:

1. Asset data → Context Collector → LLM Engine → Business Impact Modeler → Score Calculator → Reporting Engine.
2. Feedback loop: Adjust risk weights based on historical incident outcomes.

---

## 2. Component Implementation Details

### 2.1 Asset Context Collector

**Objective:** Aggregate asset metadata for risk contextualization.

**Tools & Workflow (e.g.):**

- **Data Sources:**
  - *CMDB*: ServiceNow or Jira Service Management, etc.
  - *Cloud APIs*: AWS Resource Groups, Azure Resource Graph, etc.
  - *Criticality Tags*: Custom tags (e.g., "PCI-DSS Tier 1").
- **Metadata Extraction:**
  - **Asset Type**: Classify using regex or ML (e.g., "web server" vs. "database").
  - **Exposure Level**: Check network segmentation (public IP, VLAN, etc).
  - **Criticality**: Pull from CMDB or tag systems (e.g., "high" for ERP systems).
- **Example Code [python]:**

```
def get_asset_context(asset_id):  
    cmdb_data = servicenow_api.get_asset(asset_id)  
    cloud_data = aws_api.describe_instance(asset_id)  
    return {  
        "type": cmdb_data["type"],  
        "exposure": "public" if cloud_data["public_ip"] else "internal",  
        "criticality": cmdb_data["tags"].get("criticality", "medium")  
    }
```

---

### 2.2 LLM-Based Risk Reasoning Engine

**Objective:** Generate contextual risk narratives using LLMs.

**Implementation Steps (e.g.):**

1. **LLM Setup:**
  - *Base Model*: GPT-4 or fine-tuned Llama 2 on cybersecurity datasets (CVE descriptions, MITRE ATT&CK, etc).

- *Prompt Engineering:*

```
prompt = f"""
Vulnerability: {cve_description}
Asset: {asset_type}, {exposure}, {criticality}.
Explain the risk in business terms.
"""

response = openai.ChatCompletion.create(model="gpt-4", messages=[{"role": "user", "content": prompt}])
# Output: "Remote code execution on a public-facing payment gateway →
High risk of data breach and PCI-DSS non-compliance."
```

## 2. Validation:

- Compare LLM outputs with expert assessments (precision/recall, etc).

---

## 2.3 Business Impact Modeler

**Objective:** Quantify operational, financial, and reputational risks.

**Implementation Strategies (e.g.):**

- **Cost Models:**

- *Downtime Costs:*  $(\text{hourly revenue}) \times (\text{MTTR estimate})$ .
- *Regulatory Fines:* Use GDPR calculator (4% of global turnover) or HIPAA penalties, etc.
- *Reputational Damage:* Estimate via historical breach data (e.g., 20% customer churn).

- **Example Code:**

```
def calculate_fines(asset_type, region):
    if "GDPR" in asset_type["compliance_tags"]:
        return revenue * 0.04
    elif "HIPAA" in asset_type["compliance_tags"]:
        return 50000 # Base penalty
```

## 2.4 Prioritization Score Calculator

**Objective:** Compute dynamic risk scores.

**Implementation Steps (e.g.):**

1. **Weighted Formula:**

```
Risk Score = (LLM Threat Level × 0.5) + (Business Impact × 0.3) + (Exposure × 0.2)
```

- *LLM Threat Level*: Normalized score from 1–10 based on risk narrative.
- *Business Impact*: Financial/reputational impact (Euro).
- *Exposure*: Public (1.0) vs internal (0.5).

2. **Example Code:**

```
def calculate_score(threat_level, business_impact, exposure):  
    return (threat_level * 0.5) + (business_impact * 0.3) + (exposure * 0.2)
```

---

## 2.5 Reporting Engine

**Objective:** Generate actionable reports for stakeholders.

**Implementation Details (e.g.):**

- **Dashboard:**
  - *Tools*: Grafana or Tableau for visualizations, etc.
  - *Metrics*: Top vulnerabilities by risk score, cost projections, etc.
- **Remediation Playbooks:**
  - Auto-generate Jira tickets with priority labels (e.g., "Critical: Patch CVE-2023-XXXX within 48h").
- **Example Output:**

```
## Priority 1: CVE-2023-XXXX (Score: 9.8/10)  
- Asset: Public-facing payment gateway (High criticality).  
- Business Impact: $2M potential downtime + $1.2M GDPR fines.  
- Action: Apply patch KB123456 by [date].
```

---



### 3. Evaluation Metrics

1. **Accuracy:**
    - Compare LLM risk narratives vs expert assessments (F1-score).
  2. **Business Alignment:**
    - Reduction in downtime costs post-remediation (e.g., 30% faster patching).
  3. **Usability:**
    - Stakeholder feedback on report clarity (Likert scale).
- 

### 4. Challenges & Mitigation (optional)

- **Data Silos:**
    - Use middleware (Apache NiFi) to unify CMDB/cloud data.
  - **LLM Hallucinations:**
    - Rule-based validation (e.g., flag non-existent CVEs).
  - **Dynamic Environments:**
    - Schedule periodic asset context refreshes (e.g., nightly scans).
- 

### 5. Tools & Technologies (e.g.)

- **LLM Framework:** Hugging Face Transformers, OpenAI API, Claude ai, DeepSeek, etc.
  - **Data Integration:** Apache NiFi, AWS SDK, etc.
  - **Visualization:** Grafana, Tableau, etc.
  - **Database:** PostgreSQL, MongoDB, etc.
-