

# Approccio SWARM: tool VeriSmart


Salvatore La Torre



*Dipartimento di Informatica  
Università degli Studi di Salerno*

# Concurrency makes bug finding harder.


**State space explosion** (i.e., large number of interleavings):



Problem: modern hardware means concurrency is everywhere  
⇒ software is increasingly concurrent


# Concurrency makes bug finding *easier*.

Concurrent hardware allows us to **run many (smaller) analysis tasks in parallel**:



# Concurrency makes bug finding *easier*.

Concurrent hardware allows us to **run many (smaller) analysis tasks in parallel**:



How can we partition a task into **independent smaller** tasks?

# Strategy competition vs. task competition

Strategy competition: run **different settings** on **same task**  
(**first counterexample** “wins” and aborts other tasks)

## Swarm Verification Techniques

Gerard J. Holzmann, Rajeev Joshi, and Alex Groce

“anticipate the appearance of systems with large numbers of CPU cores, but without matching increases in clockspeeds... describe a model checking strategy that leverages this trend”

Abs the ben in processing speeds, the... makers started redirecting their efforts to the development of multi-core systems. For the near-term future, we can anticipate the appearance of systems with large numbers of CPU cores, but without matching increases in clockspeeds. We will describe a model checking strategy that can allow us to leverage this trend, and that allows us to tackle significantly larger problem sizes than before.

**Index Terms**—software engineering tools and techniques, logic model checking, distributed algorithms, software verification.




Task competition: run **same prover** (setting) on **different tasks**  
⇒ How can we partition a task into **independent smaller** tasks?

# Reduced interleaving instances

Our goal:

Split set of interleavings  $I_k(P)$  into **subsets** that can be **analyzed symbolically** and **independently**.




Our solution:

Derive program **variants**  $P_g$  that **allow context switches only** in **subsets** of statements (**tiles**) s.t.  $I_k(P) = \bigcup_g I_k(P_g)$ .

# Reduced interleaving instances

Goal:

**Split** set of interleavings  $I_k(P)$  into **subsets** that can be **analyzed symbolically** and **independently**.



Idea:

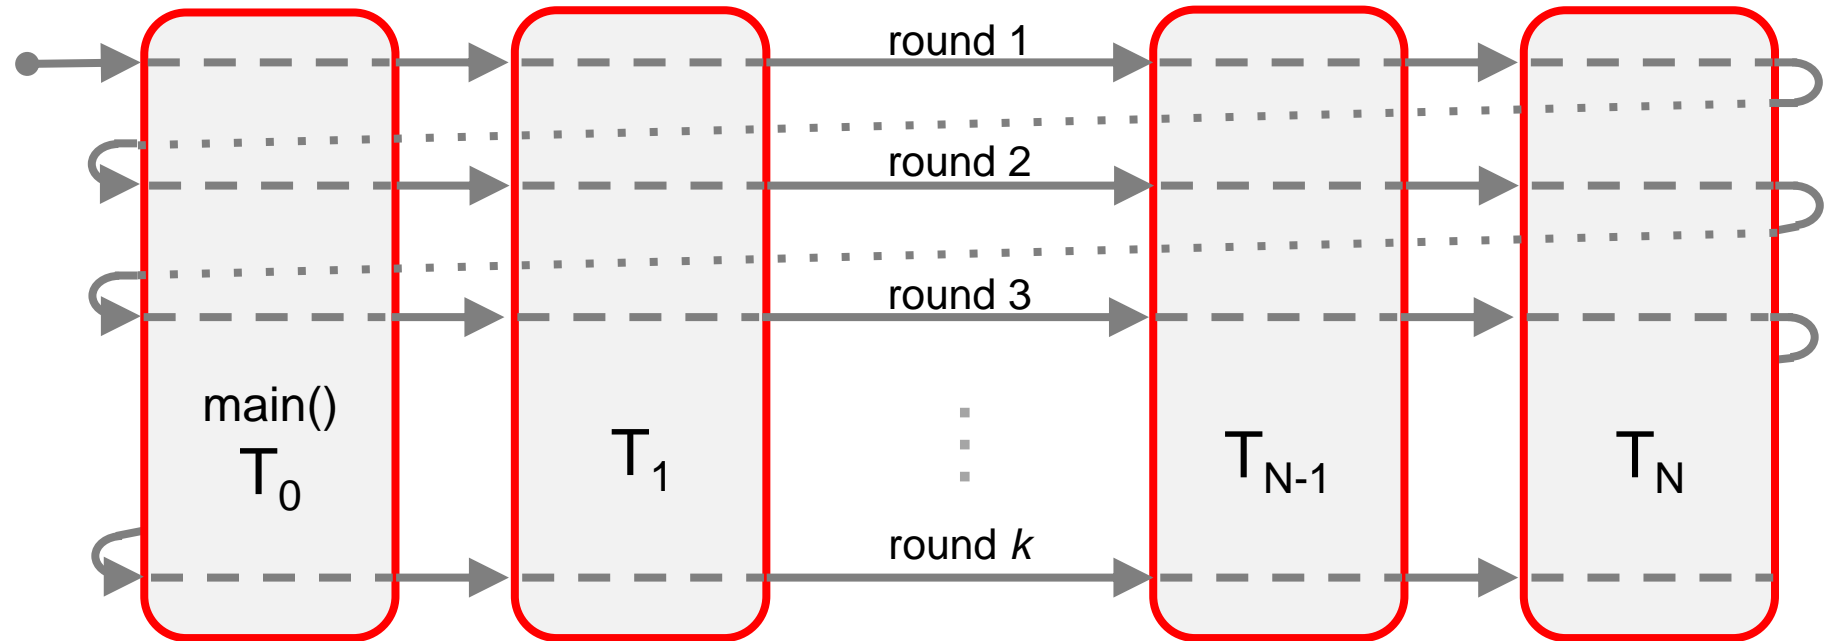
**Derive** program **variants**  $P_{\mathcal{I}}$  that **allow context switches only** in **subsets** of statements (**tiles**) s.t.  $I_k(P) = \bigcup_{\mathcal{I}} I_k(P_{\mathcal{I}})$ .

# Tiling Threads



# Tiling threads


## Assumption: bounded concurrent programs



- finite #threads, fixed (but arbitrary) schedule
  - captures all bounded round-robin computations for given bound
- bugs manifest within very few rounds [Musuvathi, Qadeer, PLDI'07]

# Tiling threads


**Assumption: bounded concurrent programs**



- finite #stmts
- control can only go forward
  - simplifies analysis and tiling

# Tiling threads

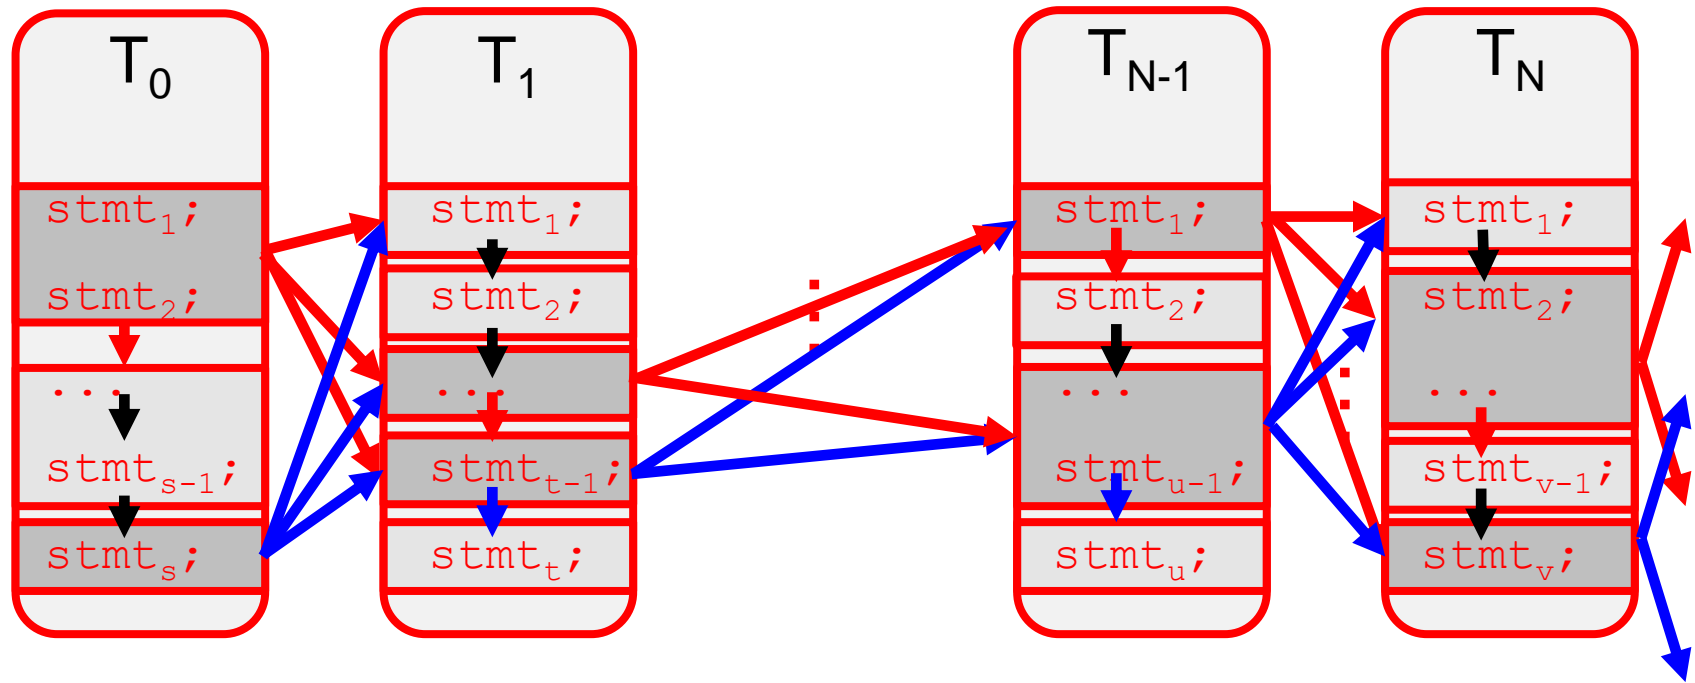
Tiles:



- **tile**: (contiguous) subset of visible statements
  - other tile types possible: random subsets, data-flow driven, ...
- **tiling**: partition of program into tiles
- **uniform window tiling**: all tiles have same size
  - number of visible statements

# Tiling threads

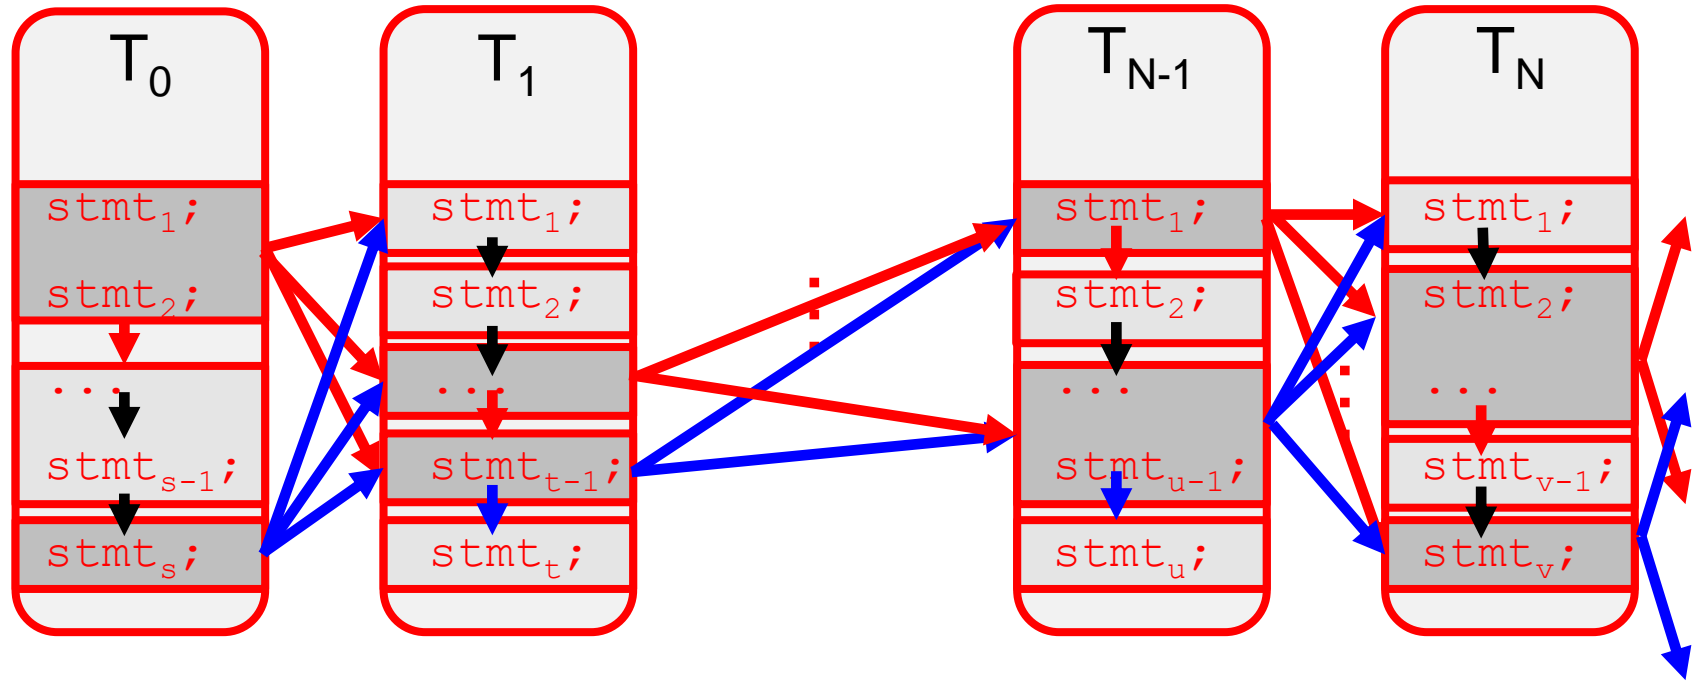
## Tile selection:



- **z-selection**: subset of  $z$  tiles for each thread
  - **context switches** are **only** allowed from **selected tiles**  
 $\Rightarrow$  context switches can only go into other selected tiles (or first thread statement)
- each z-selection specifies a **reduced interleaving instance**

# Tiling threads

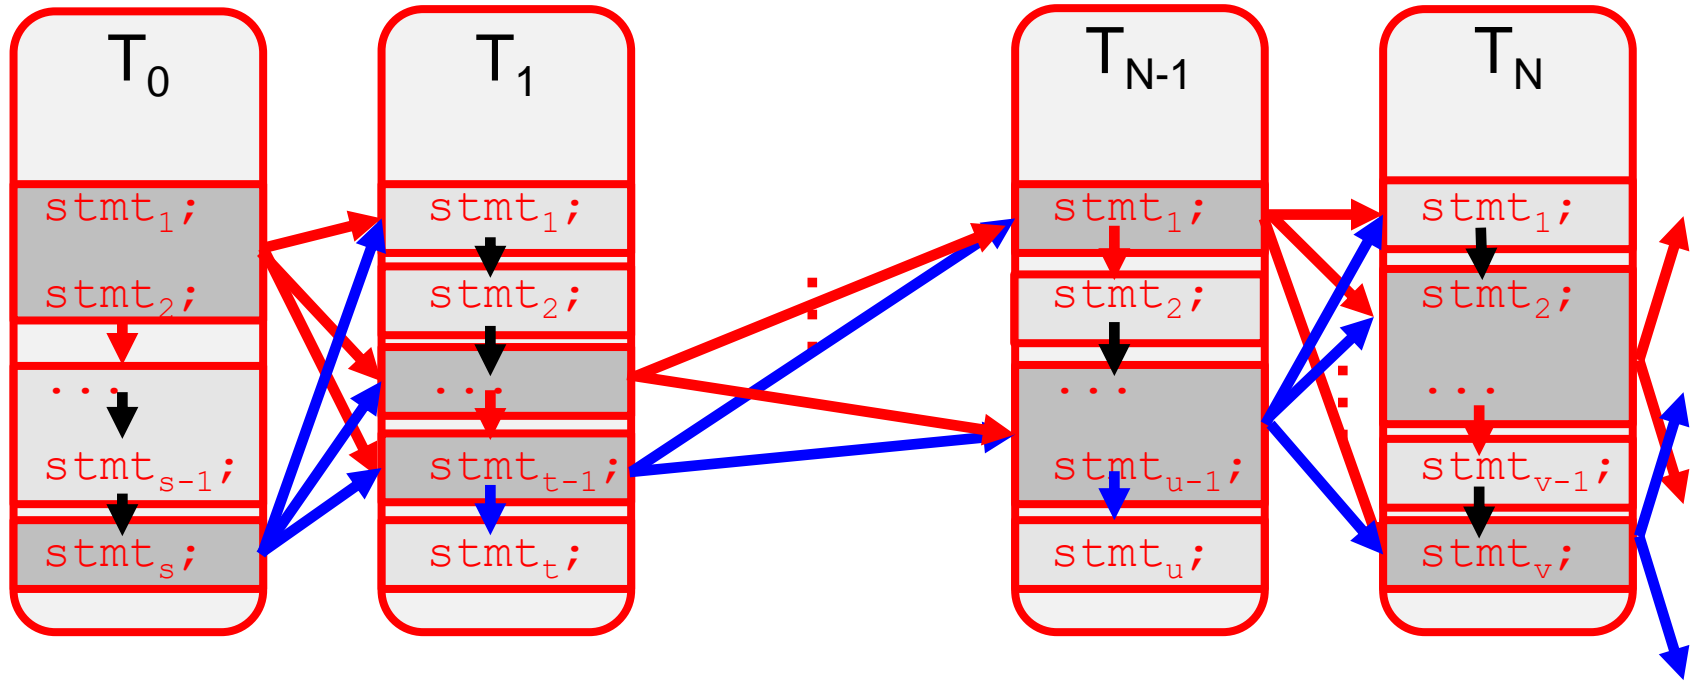
## Completeness of selections:



- each interleaving with  $k$  context switches can be covered by a  $\lceil k/2 \rceil$ -selection  $\vartheta \in \Theta_P$ 
    - each thread can only switch out at most  $\lceil k/2 \rceil$  times
- $\Rightarrow$  set of all  $\lceil k/2 \rceil$ -selections together covers all interleavings with  $k$  context switches:  $I_k(P) = \bigcup_{\vartheta} I_k(P_{\vartheta})$

# Tiling threads

## Completeness of selections:



- number of selections grows exponentially  
⇒ sampling

# **VERISmart**

**(Verification Smart)**

# VERISmart implements swarm verification by task competition for multi-threaded C.

Target:

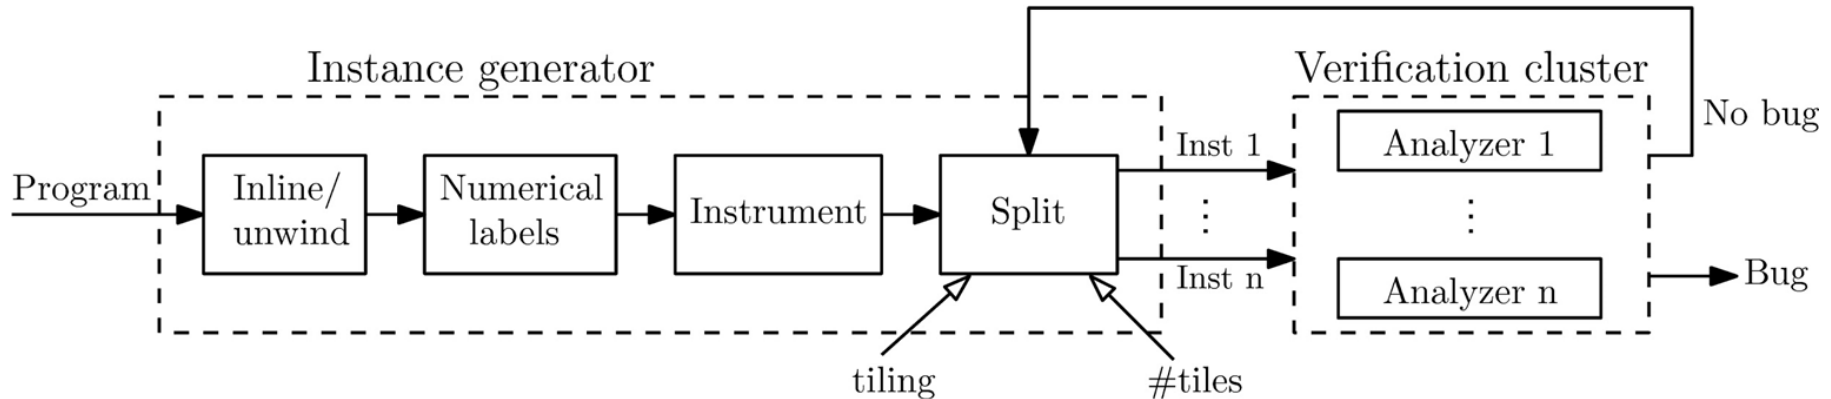
- C programs with “**rare**” concurrency bugs, i.e.,
  - “**large**” number of interleavings
  - “**few**” interleavings lead to a bug
- **automatic bug-finding** (bounded analysis, not complete)
- **reachability**
  - assertion failure
  - out-of-bound array, division-by-zero, ...
  - linearizability → reachability [Bouajjani et al., POPL’15, ICALP’15]

Approach:

- source-to-source translation to generate instances (for tiling)
  - instances are bounded concurrent programs
- use cluster to run Lazy-CSeq over instances [Inverso et al., CAV’14]



# VERISmart architecture



- **Inline/unwind** module:
  - concurrent program → bounded concurrent program
- **Numerical labels** module:
  - inject numerical labels at each visible statement
- **Instrument** module:
  - instrument the code with guarded commands (**yield**) that can enable/disable context switch points at numerical labels
- **Split** module:
  - generate variants with configuration from tiling and #tiles
  - randomize number of generated variants when #variants is large

# Why does this work?

Remember:

Each  $P_{\vartheta}$  allows only a (small) subset of  $P$ 's interleavings

We assume bugs are rare,

- so for most  $\vartheta$ ,  $P_{\vartheta}$  does not exhibit the bug...
- ... and the analysis will run out of time
- but if  $P_{\vartheta}$  does exhibit the bug...
- ... the analysis will find it quick(er)

Hence,

- overall CPU time consumption goes (way) up...
- ... but with enough cores CPU time is free and...
- mean wall clock time to find failure goes down

# **Experimental Evaluation**

# ABA problem

- Concurrency problem that can occur when a thread reads a shared variable twice and another thread accesses it between the two reads
- Sequence of events that identify an ABA problem:
  - Thread 1 reads A from shared memory, and then is preempted
  - Thread 2 modifies this value from A to B and then back to A
  - Thread 1 is resumed and perceives that nothing has changed

# Why is ABA a problem?

- Thread 1 acts as nothing has changed but ...  
..... this is not the case and can lead to error
- For example, in lock-free data structures:
  - an item X is removed from the structure
  - a new item Y is added to the structure
  - due to optimizations, it is common that Y is allocated at the same location as X
  - pointer to Y is same as pointer to X
  - however Y might be different from X

# Experiments on lock-free data structures

eliminationstack:

- ABA problem: requires 7 threads for exposure
- Lazy-CSeq can find bug in **~13h** and **4GB**
  - #unwind=1, #rounds=2, #threads=8, #visible=52
- all other tools fail

safestack:

- ABA problem: requires context bound of 5
- Lazy-CSeq can find bug in **~7h** and **6.5GB**
  - #unwind=3, #rounds=4, #threads=4, #visible=152
- all other tools fail

# eliminationstack: Results

- Lazy-CSeq: 46764 sec, 4.2 GB
- CBMC (sequential): 80.8 sec, 0.7 GB
  - average over 3000 interleavings, bug not found

| VERISmart: 2 tiles           |        |        | VERISmart: 2 tiles           |        |        | VERISmart: 2 tiles           |         |         |
|------------------------------|--------|--------|------------------------------|--------|--------|------------------------------|---------|---------|
| #1: tile size 12, t_max 3hrs |        |        | #2: tile size 14, t_max 3hrs |        |        | #3: tile size 18, t_max 3hrs |         |         |
| Verification                 | Time   | Memory | Verification                 | Time   | Memory | Verification                 | Time    | Memory  |
| Min                          | 34.9   | 945.2  | Min                          | 39.7   | 979.84 | Min                          | 37.1    | 999.8   |
| Max                          | 4753.6 | 1199.1 | Max                          | 7195.2 | 1281.3 | Max                          | 10762.0 | 1785.5  |
| Average                      | 1116.3 | 1017.8 | Average                      | 2169.5 | 1096.3 | Average                      | 3162.41 | 1156.91 |
| instances with bug: 38.33%   |        |        | instances with bug: 61.38%   |        |        | instances with bug: 69.01%   |         |         |

fastest instances very fast – 1000x

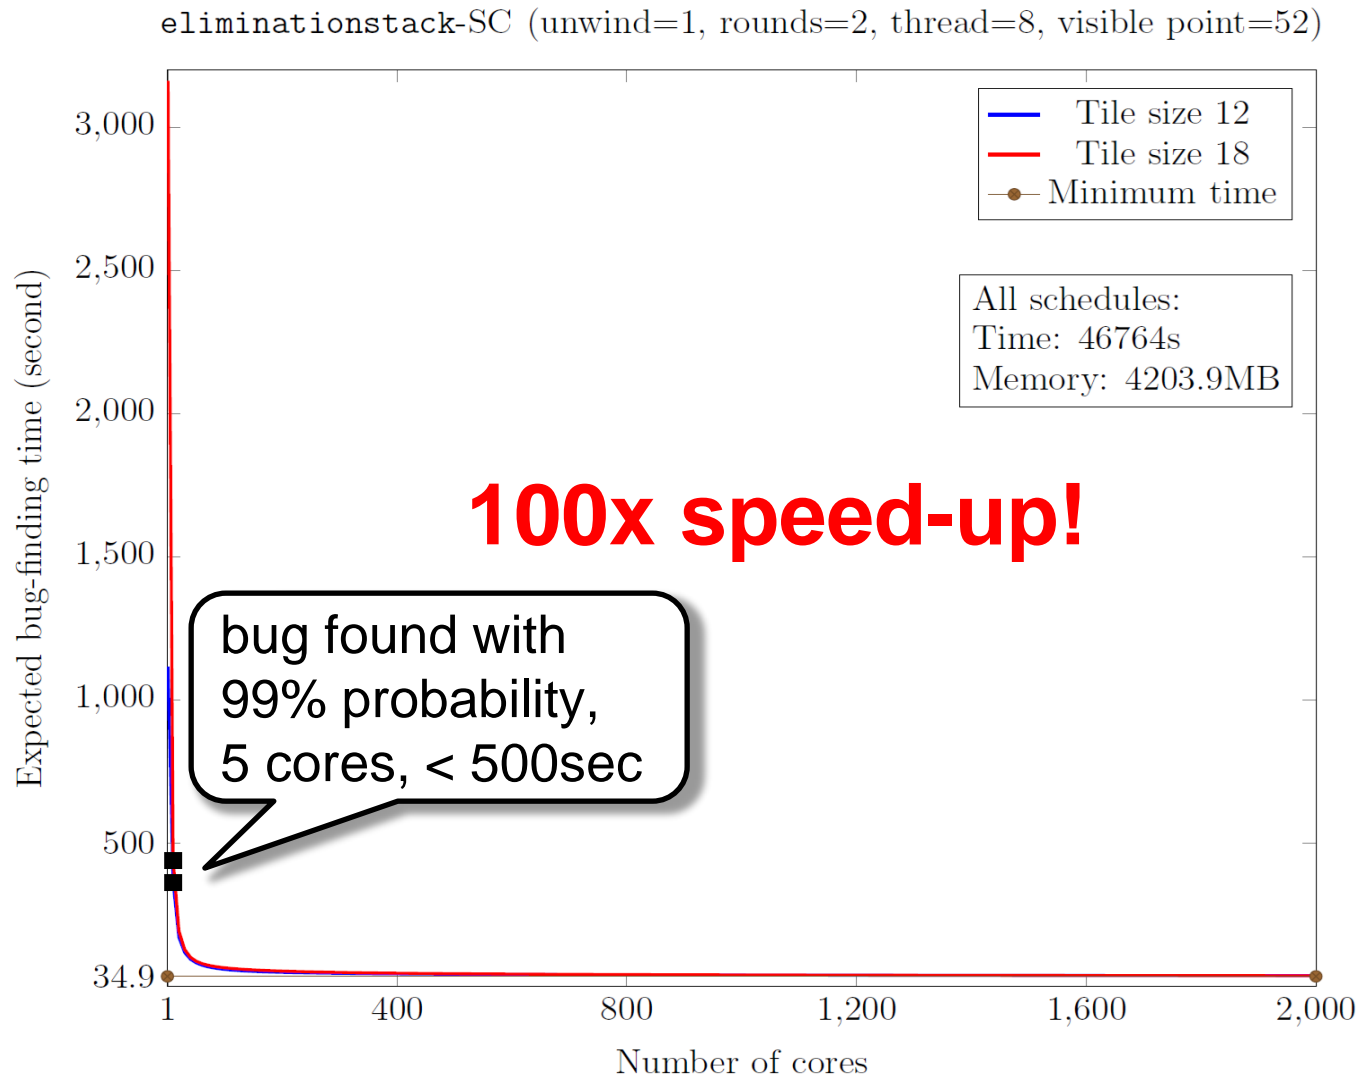
reduced memory consumption – 4x

average still very fast – 40x

high fraction of bug-exposing instances

some slowdown for larger tile sizes – 10x

# eliminationstack: Expected bug finding time





# safestack (SC): Results

- Lazy-CSeq: 24139 sec, 6.6 GB
- CBMC (sequential): 55.4 sec, 0.7 GB
  - average over 3000 interleavings, bug not found

VERISmart: 4 tiles per thread


| #1: tile size 11, t_max 1hr |        |        | #2: tile size 14, t_max 1hr |        |        | #3: tile size 20, t_max 4hrs |         |        |
|-----------------------------|--------|--------|-----------------------------|--------|--------|------------------------------|---------|--------|
| Verification                | Time   | Memory | Verification                | Time   | Memory | Verification                 | Time    | Memory |
| Min                         | 195.6  | 774.5  | Min                         | 574.8  | 846.6  | Min                          | 313.0   | 850.3  |
| Max                         | 2662.6 | 1265.7 | Max                         | 3521.8 | 1450.4 | Max                          | 10315.8 | 3830.8 |
| Average                     | 1172.2 | 928.8  | Average                     | 1851.1 | 1147.3 | Average                      | 2167.5  | 1230.1 |
| instances with bug: 1.26%   |        |        | instances with bug: 2.14%   |        |        | instances with bug: 10.20%   |         |        |

lower fraction of bug-exposing instances than eliminationstack

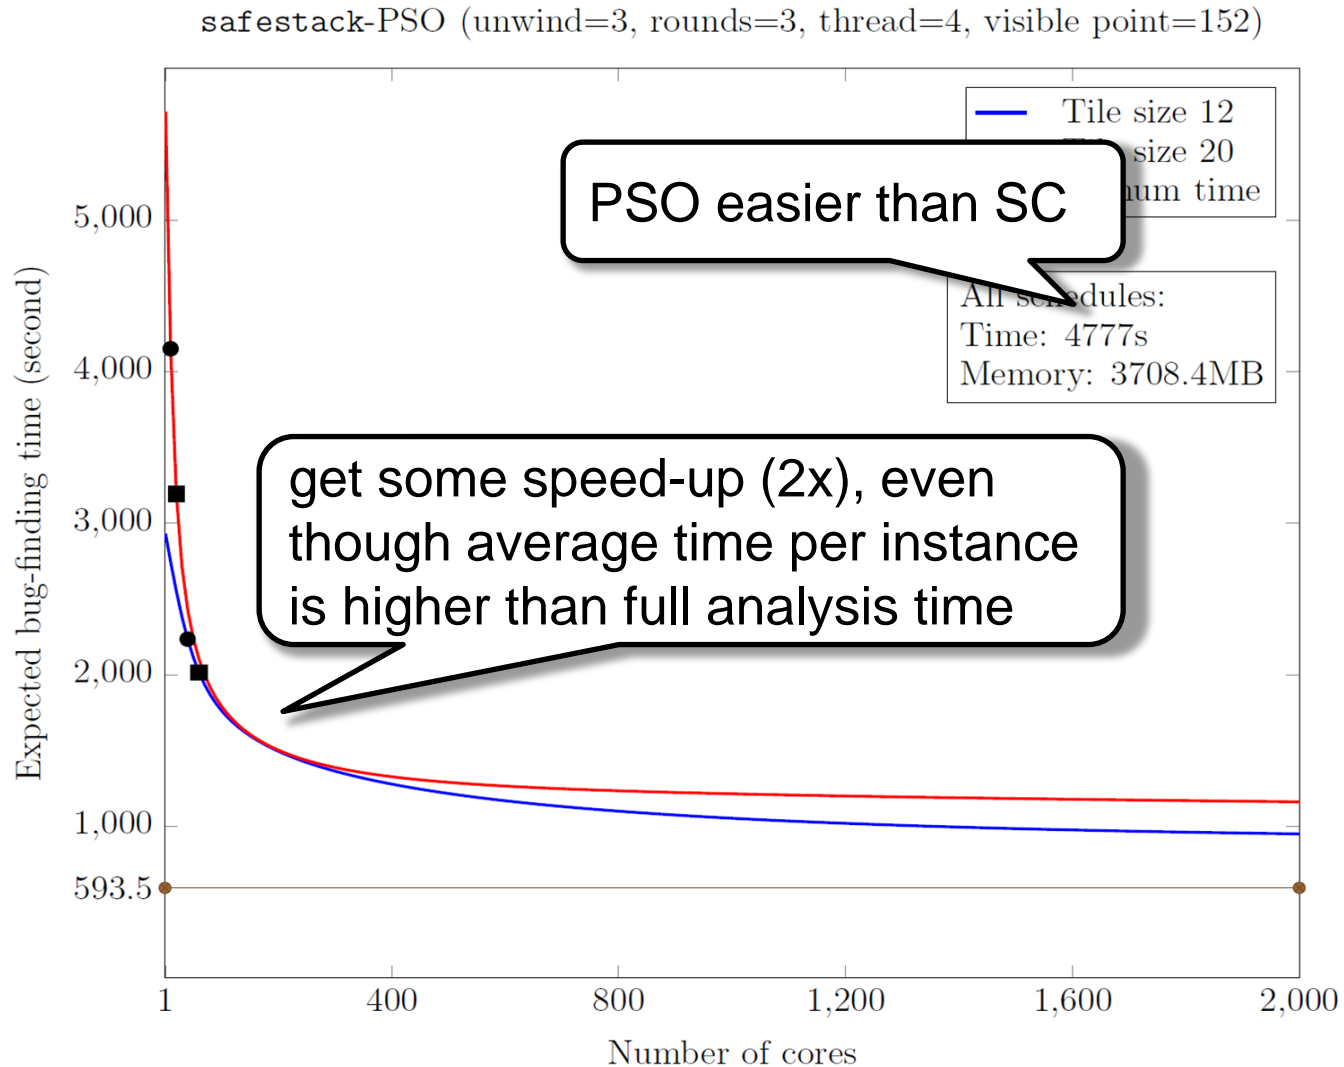
...but boosted with larger tile sizes

⇒ similar picture, but less advantage for VERISmart

# safestack (SC): Expected bug finding time



# safestack (PSO): Expected bug finding time



# Conclusions

- first task-competitive swarm verification approach
- exploits availability of many cores to reduce mean wall clock time to find failure
  - allows us to handle very hard problems
  - high speed-ups already for 5-50 cores
- reduced interleaving instances boost bug-finding capabilities

# Future Work

- production-quality implementation based on LLVM
- other backends (testing)
- other tiling styles
- fast over-approximations to filter out safe instances

# References

- Parallel Bug-finding in Concurrent Programs via Reduced Interleaving Instances  
[Nguyen-Schrammel-Fischer-La Torre-Parlato, ASE'17]
- Swarm verification techniques  
[Holzmann-Joshi-Groce, IEEE Trans. Software Eng., 2011]
- Swarm verification [Holzmann-Joshi-Groce, ASE'08]
- Cloud-based verification of concurrent software  
[Holzmann, VMCAI'16]